

wine_quality_red_and_white

May 11, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: import warnings
warnings.filterwarnings('ignore')
```

```
[3]: red_wine = pd.read_csv("winequality-red.csv")
white_wine = pd.read_csv("winequality-white.csv")
```

```
[4]: red_wine.shape
```

```
[4]: (1599, 12)
```

```
[5]: white_wine.shape
```

```
[5]: (4898, 12)
```

```
[6]: wine = pd.concat([red_wine, white_wine], ignore_index=True)
```

```
[7]: wine_df = pd.concat([red_wine, white_wine], ignore_index=True)
```

```
[8]: wine.shape
```

```
[8]: (6497, 12)
```

```
[9]: wine.head()
```

```
[9]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4           0.70         0.00           1.9       0.076 \
1           7.8           0.88         0.00           2.6       0.098
2           7.8           0.76         0.04           2.3       0.092
3          11.2           0.28         0.56           1.9       0.075
4           7.4           0.70         0.00           1.9       0.076

   free sulfur dioxide  total sulfur dioxide  density    pH  sulphates
0             11.0           34.0    0.9978  3.51       0.56 \
```

1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[10]: wine.tail()
```

```
[10]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
6492           6.2             0.21         0.29           1.6       0.039 \
6493           6.6             0.32         0.36           8.0       0.047
6494           6.5             0.24         0.19           1.2       0.041
6495           5.5             0.29         0.30           1.1       0.022
6496           6.0             0.21         0.38           0.8       0.020
```

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
6492	24.0	92.0	0.99114	3.27	0.50 \
6493	57.0	168.0	0.99490	3.15	0.46
6494	30.0	111.0	0.99254	2.99	0.46
6495	20.0	110.0	0.98869	3.34	0.38
6496	22.0	98.0	0.98941	3.26	0.32

	alcohol	quality
6492	11.2	6
6493	9.6	5
6494	9.4	6
6495	12.8	7
6496	11.8	6

```
[11]: wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6497 entries, 0 to 6496
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	fixed acidity	6497 non-null	float64
1	volatile acidity	6497 non-null	float64
2	citric acid	6497 non-null	float64
3	residual sugar	6497 non-null	float64
4	chlorides	6497 non-null	float64
5	free sulfur dioxide	6497 non-null	float64

```

6   total sulfur dioxide  6497 non-null  float64
7   density              6497 non-null  float64
8   pH                  6497 non-null  float64
9   sulphates            6497 non-null  float64
10  alcohol              6497 non-null  float64
11  quality              6497 non-null  int64
dtypes: float64(11), int64(1)
memory usage: 609.2 KB

```

```
[12]: wine.isnull().sum()
```

```

[12]: fixed acidity      0
      volatile acidity   0
      citric acid        0
      residual sugar     0
      chlorides          0
      free sulfur dioxide 0
      total sulfur dioxide 0
      density            0
      pH                0
      sulphates          0
      alcohol            0
      quality            0
      dtype: int64

```

```
[13]: wine.describe()
```

```

[13]:
      fixed acidity  volatile acidity  citric acid  residual sugar
count    6497.000000      6497.000000  6497.000000    6497.000000 \
mean         7.215307         0.339666    0.318633         5.443235
std         1.296434         0.164636    0.145318         4.757804
min          3.800000         0.080000    0.000000         0.600000
25%          6.400000         0.230000    0.250000         1.800000
50%          7.000000         0.290000    0.310000         3.000000
75%          7.700000         0.400000    0.390000         8.100000
max         15.900000         1.580000    1.660000        65.800000

      chlorides  free sulfur dioxide  total sulfur dioxide      density
count    6497.000000      6497.000000    6497.000000  6497.000000 \
mean         0.056034        30.525319    115.744574    0.994697
std         0.035034        17.749400     56.521855    0.002999
min          0.009000         1.000000         6.000000    0.987110
25%          0.038000        17.000000        77.000000    0.992340
50%          0.047000        29.000000       118.000000    0.994890
75%          0.065000        41.000000       156.000000    0.996990
max          0.611000       289.000000       440.000000    1.038980

      pH  sulphates  alcohol  quality

```


count	6497.000000	6497.000000	6497.000000	6497.000000
mean	3.218501	0.531268	10.491801	5.818378
std	0.160787	0.148806	1.192712	0.873255
min	2.720000	0.220000	8.000000	3.000000
25%	3.110000	0.430000	9.500000	5.000000
50%	3.210000	0.510000	10.300000	6.000000
75%	3.320000	0.600000	11.300000	6.000000
max	4.010000	2.000000	14.900000	9.000000

0.1 Outlier detection and Visualization

```
[14]: from scipy import stats
import seaborn as sns
```

```
[15]: plt.subplots(figsize=(15, 9))
plt.xticks(rotation=45)
sns.boxplot(data=wine)
```

[15]: <Axes: >




```
[16]: # calculate z-scores and identify outliers
z_scores = stats.zscore(wine['total sulfur dioxide'])
threshold = 3
```

```

outliers = wine[np.abs(z_scores) > threshold]

# create scatter plot
fig, ax = plt.subplots(figsize=(15, 9))
ax.scatter(wine.index, wine['total sulfur dioxide'], alpha=0.5)
ax.scatter(outliers.index, outliers['total sulfur dioxide'], color='r')
plt.show()

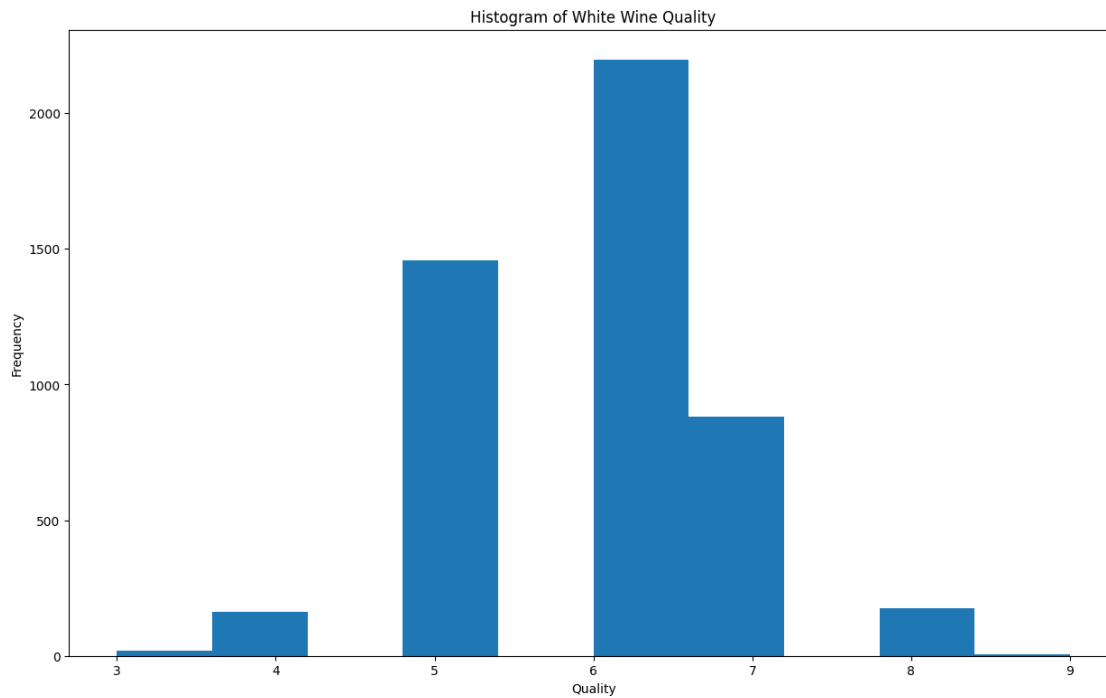
```



```

[17]: # create a histogram with larger figure size
fig, ax = plt.subplots(figsize=(15, 9))
ax.hist(white_wine['quality'])
ax.set_xlabel('Quality')
ax.set_ylabel('Frequency')
ax.set_title('Histogram of White Wine Quality')
plt.show()


```



```
[18]: # create scatter plot with larger figure size
ig, ax = plt.subplots(figsize=(15, 9))
ax.scatter(x='alcohol', y='quality', data=white_wine, alpha=0.5)

# set style property
ax.set(xlabel='Alcohol', ylabel='Quality')
plt.style.use('seaborn') # set style to seaborn


plt.show()
```



```
[19]: fig, ax = plt.subplots(figsize=(15, 9))
ax.boxplot([white_wine['quality'], white_wine['alcohol']], labels=['Quality',
↪ 'Alcohol'])

# set style property
ax.set(xlabel='Variable', ylabel='Value')
plt.style.use('seaborn')

plt.show()
```



0.1.1 Splitting the data

```
[20]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
```

```
[21]: X = wine.drop('quality', axis=1) # Extract the input features
      y = wine['quality'] # Extract the target variable
```

```
[22]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      random_state=42)
```

```
[23]: sc = StandardScaler()
```

```
[24]: X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

0.2 1. Linear regression model

```
[25]: from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

```
[26]: # Create a linear regression model
      lr_model = LinearRegression()
```



```
[27]: # Fit the model on the training set
lr_model.fit(X_train, y_train)
```

```
[27]: LinearRegression()
```

```
[28]: # Make predictions on the testing set
y_pred = lr_model.predict(X_test)
```

```
[29]: # Calculate RMSE, MAE, MSE, and R2 score
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
[30]: print('RMSE:', rmse)
print('MAE:', mae)
print('MSE:', mse)
print('R-squared Score:', r2)
```

RMSE: 0.7393892357611412

MAE: 0.5658710079723465

MSE: 0.5466964419594444

R-squared Score: 0.2597673129771396

0.3 2. XGBoost

```
[31]: from xgboost import XGBRegressor
```

```
[32]: # Create an XGBoost model
xgb_model = XGBRegressor(objective='reg:squarederror', random_state=42)
```

```
[33]: xgb_model.fit(X_train, y_train)
```

```
[33]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                 colsample_bylevel=None, colsample_bynode=None,
                 colsample_bytree=None, early_stopping_rounds=None,
                 enable_categorical=False, eval_metric=None, feature_types=None,
                 gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
                 interaction_constraints=None, learning_rate=None, max_bin=None,
                 max_cat_threshold=None, max_cat_to_onehot=None,
                 max_delta_step=None, max_depth=None, max_leaves=None,
                 min_child_weight=None, missing=nan, monotone_constraints=None,
                 n_estimators=100, n_jobs=None, num_parallel_tree=None,
                 predictor=None, random_state=42, ...)
```

```
[34]: # Make predictions on the testing set
y_pred = xgb_model.predict(X_test)
```

```
[35]: # Calculate RMSE, MAE, MSE, and R2 score
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
[36]: print('RMSE:', rmse)
print('MAE:', mae)
print('MSE:', mse)
print('R-squared Score:', r2)
```

```
RMSE: 0.6290334979171012
MAE: 0.46276961509998027
MSE: 0.39568314150182365
R-squared Score: 0.4642408975742527
```

0.4 3. ANN

```
[37]: import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam
```

```
[38]: # Define hyperparameters
learning_rate = 0.001
num_epochs = 100
batch_size = 32
```

```
[39]: # Create a neural network model
ann_model = Sequential()
ann_model.add(Dense(128, input_dim=X_train.shape[1], activation='relu'))
ann_model.add(Dense(64, activation='relu'))
ann_model.add(Dense(32, activation='relu'))
ann_model.add(Dense(16, activation='relu'))
ann_model.add(Dense(1))
```

```
[40]: # Compile model with Adam optimizer and custom learning rate
opt = Adam(lr=learning_rate)
ann_model.compile(loss='mean_squared_error', optimizer=opt,
    ↪metrics=['mean_squared_error'])
```

```
[41]: # Train model with specified batch size and number of epochs
history = ann_model.fit(X_train, y_train, batch_size=batch_size,
    ↪epochs=num_epochs, verbose=1, validation_split=0.2)
```

```
Epoch 1/100
130/130 [=====] - 1s 4ms/step - loss: 6.2541 -
mean_squared_error: 6.2541 - val_loss: 1.5355 - val_mean_squared_error: 1.5355
Epoch 2/100
```

```

130/130 [=====] - 0s 2ms/step - loss: 1.2199 -
mean_squared_error: 1.2199 - val_loss: 1.1021 - val_mean_squared_error: 1.1021
Epoch 3/100
130/130 [=====] - 0s 2ms/step - loss: 0.8532 -
mean_squared_error: 0.8532 - val_loss: 0.8042 - val_mean_squared_error: 0.8042
Epoch 4/100
130/130 [=====] - 0s 2ms/step - loss: 0.6464 -
mean_squared_error: 0.6464 - val_loss: 0.6210 - val_mean_squared_error: 0.6210
Epoch 5/100
130/130 [=====] - 0s 2ms/step - loss: 0.5498 -
mean_squared_error: 0.5498 - val_loss: 0.5433 - val_mean_squared_error: 0.5433
Epoch 6/100
130/130 [=====] - 0s 2ms/step - loss: 0.5227 -
mean_squared_error: 0.5227 - val_loss: 0.5158 - val_mean_squared_error: 0.5158
Epoch 7/100
130/130 [=====] - 0s 2ms/step - loss: 0.4941 -
mean_squared_error: 0.4941 - val_loss: 0.4747 - val_mean_squared_error: 0.4747
Epoch 8/100
130/130 [=====] - 0s 2ms/step - loss: 0.4894 -
mean_squared_error: 0.4894 - val_loss: 0.5111 - val_mean_squared_error: 0.5111
Epoch 9/100
130/130 [=====] - 0s 2ms/step - loss: 0.4689 -
mean_squared_error: 0.4689 - val_loss: 0.4761 - val_mean_squared_error: 0.4761
Epoch 10/100
130/130 [=====] - 0s 2ms/step - loss: 0.4790 -
mean_squared_error: 0.4790 - val_loss: 0.4642 - val_mean_squared_error: 0.4642
Epoch 11/100
130/130 [=====] - 0s 2ms/step - loss: 0.4620 -
mean_squared_error: 0.4620 - val_loss: 0.4596 - val_mean_squared_error: 0.4596
Epoch 12/100
130/130 [=====] - 0s 2ms/step - loss: 0.4494 -
mean_squared_error: 0.4494 - val_loss: 0.4653 - val_mean_squared_error: 0.4653
Epoch 13/100
130/130 [=====] - 0s 2ms/step - loss: 0.4469 -
mean_squared_error: 0.4469 - val_loss: 0.5055 - val_mean_squared_error: 0.5055
Epoch 14/100
130/130 [=====] - 0s 2ms/step - loss: 0.4445 -
mean_squared_error: 0.4445 - val_loss: 0.4848 - val_mean_squared_error: 0.4848
Epoch 15/100
130/130 [=====] - 0s 2ms/step - loss: 0.4559 -
mean_squared_error: 0.4559 - val_loss: 0.5377 - val_mean_squared_error: 0.5377
Epoch 16/100
130/130 [=====] - 0s 3ms/step - loss: 0.4333 -
mean_squared_error: 0.4333 - val_loss: 0.4706 - val_mean_squared_error: 0.4706
Epoch 17/100
130/130 [=====] - 0s 2ms/step - loss: 0.4378 -
mean_squared_error: 0.4378 - val_loss: 0.4500 - val_mean_squared_error: 0.4500
Epoch 18/100

```

```

130/130 [=====] - 0s 2ms/step - loss: 0.4318 -
mean_squared_error: 0.4318 - val_loss: 0.4828 - val_mean_squared_error: 0.4828
Epoch 19/100
130/130 [=====] - 0s 2ms/step - loss: 0.4246 -
mean_squared_error: 0.4246 - val_loss: 0.4553 - val_mean_squared_error: 0.4553
Epoch 20/100
130/130 [=====] - 0s 2ms/step - loss: 0.4245 -
mean_squared_error: 0.4245 - val_loss: 0.4597 - val_mean_squared_error: 0.4597
Epoch 21/100
130/130 [=====] - 0s 2ms/step - loss: 0.4227 -
mean_squared_error: 0.4227 - val_loss: 0.4772 - val_mean_squared_error: 0.4772
Epoch 22/100
130/130 [=====] - 0s 2ms/step - loss: 0.4118 -
mean_squared_error: 0.4118 - val_loss: 0.4584 - val_mean_squared_error: 0.4584
Epoch 23/100
130/130 [=====] - 0s 2ms/step - loss: 0.4163 -
mean_squared_error: 0.4163 - val_loss: 0.4633 - val_mean_squared_error: 0.4633
Epoch 24/100
130/130 [=====] - 0s 2ms/step - loss: 0.3978 -
mean_squared_error: 0.3978 - val_loss: 0.4825 - val_mean_squared_error: 0.4825
Epoch 25/100
130/130 [=====] - 0s 3ms/step - loss: 0.3949 -
mean_squared_error: 0.3949 - val_loss: 0.4963 - val_mean_squared_error: 0.4963
Epoch 26/100
130/130 [=====] - 0s 3ms/step - loss: 0.3999 -
mean_squared_error: 0.3999 - val_loss: 0.4645 - val_mean_squared_error: 0.4645
Epoch 27/100
130/130 [=====] - 0s 2ms/step - loss: 0.3931 -
mean_squared_error: 0.3931 - val_loss: 0.4580 - val_mean_squared_error: 0.4580
Epoch 28/100
130/130 [=====] - 0s 2ms/step - loss: 0.3891 -
mean_squared_error: 0.3891 - val_loss: 0.4778 - val_mean_squared_error: 0.4778
Epoch 29/100
130/130 [=====] - 0s 2ms/step - loss: 0.3856 -
mean_squared_error: 0.3856 - val_loss: 0.4630 - val_mean_squared_error: 0.4630
Epoch 30/100
130/130 [=====] - 0s 2ms/step - loss: 0.3777 -
mean_squared_error: 0.3777 - val_loss: 0.4740 - val_mean_squared_error: 0.4740
Epoch 31/100
130/130 [=====] - 0s 2ms/step - loss: 0.3792 -
mean_squared_error: 0.3792 - val_loss: 0.4703 - val_mean_squared_error: 0.4703
Epoch 32/100
130/130 [=====] - 0s 2ms/step - loss: 0.3690 -
mean_squared_error: 0.3690 - val_loss: 0.4700 - val_mean_squared_error: 0.4700
Epoch 33/100
130/130 [=====] - 0s 2ms/step - loss: 0.3595 -
mean_squared_error: 0.3595 - val_loss: 0.4775 - val_mean_squared_error: 0.4775
Epoch 34/100

```

```

130/130 [=====] - 0s 2ms/step - loss: 0.3592 -
mean_squared_error: 0.3592 - val_loss: 0.4874 - val_mean_squared_error: 0.4874
Epoch 35/100
130/130 [=====] - 0s 3ms/step - loss: 0.3687 -
mean_squared_error: 0.3687 - val_loss: 0.4961 - val_mean_squared_error: 0.4961
Epoch 36/100
130/130 [=====] - 0s 3ms/step - loss: 0.3535 -
mean_squared_error: 0.3535 - val_loss: 0.4768 - val_mean_squared_error: 0.4768
Epoch 37/100
130/130 [=====] - 0s 2ms/step - loss: 0.3508 -
mean_squared_error: 0.3508 - val_loss: 0.5262 - val_mean_squared_error: 0.5262
Epoch 38/100
130/130 [=====] - 0s 3ms/step - loss: 0.3428 -
mean_squared_error: 0.3428 - val_loss: 0.4937 - val_mean_squared_error: 0.4937
Epoch 39/100
130/130 [=====] - 0s 3ms/step - loss: 0.3452 -
mean_squared_error: 0.3452 - val_loss: 0.5408 - val_mean_squared_error: 0.5408
Epoch 40/100
130/130 [=====] - 0s 3ms/step - loss: 0.3427 -
mean_squared_error: 0.3427 - val_loss: 0.4933 - val_mean_squared_error: 0.4933
Epoch 41/100
130/130 [=====] - 0s 2ms/step - loss: 0.3418 -
mean_squared_error: 0.3418 - val_loss: 0.5198 - val_mean_squared_error: 0.5198
Epoch 42/100
130/130 [=====] - 0s 2ms/step - loss: 0.3401 -
mean_squared_error: 0.3401 - val_loss: 0.4910 - val_mean_squared_error: 0.4910
Epoch 43/100
130/130 [=====] - 0s 2ms/step - loss: 0.3252 -
mean_squared_error: 0.3252 - val_loss: 0.4802 - val_mean_squared_error: 0.4802
Epoch 44/100
130/130 [=====] - 0s 2ms/step - loss: 0.3223 -
mean_squared_error: 0.3223 - val_loss: 0.5325 - val_mean_squared_error: 0.5325
Epoch 45/100
130/130 [=====] - 0s 2ms/step - loss: 0.3339 -
mean_squared_error: 0.3339 - val_loss: 0.4926 - val_mean_squared_error: 0.4926
Epoch 46/100
130/130 [=====] - 0s 2ms/step - loss: 0.3324 -
mean_squared_error: 0.3324 - val_loss: 0.5056 - val_mean_squared_error: 0.5056
Epoch 47/100
130/130 [=====] - 0s 2ms/step - loss: 0.3126 -
mean_squared_error: 0.3126 - val_loss: 0.5351 - val_mean_squared_error: 0.5351
Epoch 48/100
130/130 [=====] - 0s 3ms/step - loss: 0.3108 -
mean_squared_error: 0.3108 - val_loss: 0.5218 - val_mean_squared_error: 0.5218
Epoch 49/100
130/130 [=====] - 0s 2ms/step - loss: 0.3121 -
mean_squared_error: 0.3121 - val_loss: 0.5193 - val_mean_squared_error: 0.5193
Epoch 50/100

```

130/130 [=====] - 0s 2ms/step - loss: 0.3199 -
mean_squared_error: 0.3199 - val_loss: 0.5085 - val_mean_squared_error: 0.5085
Epoch 51/100

130/130 [=====] - 0s 2ms/step - loss: 0.3040 -
mean_squared_error: 0.3040 - val_loss: 0.5048 - val_mean_squared_error: 0.5048
Epoch 52/100

130/130 [=====] - 0s 2ms/step - loss: 0.2960 -
mean_squared_error: 0.2960 - val_loss: 0.5012 - val_mean_squared_error: 0.5012
Epoch 53/100

130/130 [=====] - 0s 2ms/step - loss: 0.3010 -
mean_squared_error: 0.3010 - val_loss: 0.5199 - val_mean_squared_error: 0.5199
Epoch 54/100

130/130 [=====] - 0s 2ms/step - loss: 0.2932 -
mean_squared_error: 0.2932 - val_loss: 0.5221 - val_mean_squared_error: 0.5221
Epoch 55/100

130/130 [=====] - 0s 2ms/step - loss: 0.2968 -
mean_squared_error: 0.2968 - val_loss: 0.5047 - val_mean_squared_error: 0.5047
Epoch 56/100

130/130 [=====] - 0s 2ms/step - loss: 0.2980 -
mean_squared_error: 0.2980 - val_loss: 0.4891 - val_mean_squared_error: 0.4891
Epoch 57/100

130/130 [=====] - 0s 2ms/step - loss: 0.2800 -
mean_squared_error: 0.2800 - val_loss: 0.5329 - val_mean_squared_error: 0.5329
Epoch 58/100

130/130 [=====] - 0s 2ms/step - loss: 0.2897 -
mean_squared_error: 0.2897 - val_loss: 0.5072 - val_mean_squared_error: 0.5072
Epoch 59/100

130/130 [=====] - 0s 2ms/step - loss: 0.2810 -
mean_squared_error: 0.2810 - val_loss: 0.5224 - val_mean_squared_error: 0.5224
Epoch 60/100

130/130 [=====] - 0s 2ms/step - loss: 0.2908 -
mean_squared_error: 0.2908 - val_loss: 0.5199 - val_mean_squared_error: 0.5199
Epoch 61/100

130/130 [=====] - 0s 2ms/step - loss: 0.2853 -
mean_squared_error: 0.2853 - val_loss: 0.5009 - val_mean_squared_error: 0.5009
Epoch 62/100

130/130 [=====] - 0s 3ms/step - loss: 0.2748 -
mean_squared_error: 0.2748 - val_loss: 0.5858 - val_mean_squared_error: 0.5858
Epoch 63/100

130/130 [=====] - 0s 2ms/step - loss: 0.2782 -
mean_squared_error: 0.2782 - val_loss: 0.5020 - val_mean_squared_error: 0.5020
Epoch 64/100

130/130 [=====] - 0s 2ms/step - loss: 0.2629 -
mean_squared_error: 0.2629 - val_loss: 0.5143 - val_mean_squared_error: 0.5143
Epoch 65/100

130/130 [=====] - 0s 3ms/step - loss: 0.2745 -
mean_squared_error: 0.2745 - val_loss: 0.5242 - val_mean_squared_error: 0.5242
Epoch 66/100

130/130 [=====] - 0s 2ms/step - loss: 0.2696 -
mean_squared_error: 0.2696 - val_loss: 0.5357 - val_mean_squared_error: 0.5357
Epoch 67/100

130/130 [=====] - 0s 2ms/step - loss: 0.2733 -
mean_squared_error: 0.2733 - val_loss: 0.5202 - val_mean_squared_error: 0.5202
Epoch 68/100

130/130 [=====] - 0s 2ms/step - loss: 0.2557 -
mean_squared_error: 0.2557 - val_loss: 0.5172 - val_mean_squared_error: 0.5172
Epoch 69/100

130/130 [=====] - 0s 2ms/step - loss: 0.2545 -
mean_squared_error: 0.2545 - val_loss: 0.5313 - val_mean_squared_error: 0.5313
Epoch 70/100

130/130 [=====] - 0s 2ms/step - loss: 0.2620 -
mean_squared_error: 0.2620 - val_loss: 0.4915 - val_mean_squared_error: 0.4915
Epoch 71/100

130/130 [=====] - 0s 2ms/step - loss: 0.2492 -
mean_squared_error: 0.2492 - val_loss: 0.5281 - val_mean_squared_error: 0.5281
Epoch 72/100

130/130 [=====] - 0s 2ms/step - loss: 0.2508 -
mean_squared_error: 0.2508 - val_loss: 0.5160 - val_mean_squared_error: 0.5160
Epoch 73/100

130/130 [=====] - 0s 2ms/step - loss: 0.2527 -
mean_squared_error: 0.2527 - val_loss: 0.5283 - val_mean_squared_error: 0.5283
Epoch 74/100

130/130 [=====] - 0s 2ms/step - loss: 0.2354 -
mean_squared_error: 0.2354 - val_loss: 0.5270 - val_mean_squared_error: 0.5270
Epoch 75/100

130/130 [=====] - 0s 2ms/step - loss: 0.2463 -
mean_squared_error: 0.2463 - val_loss: 0.5150 - val_mean_squared_error: 0.5150
Epoch 76/100

130/130 [=====] - 0s 2ms/step - loss: 0.2412 -
mean_squared_error: 0.2412 - val_loss: 0.4939 - val_mean_squared_error: 0.4939
Epoch 77/100

130/130 [=====] - 0s 2ms/step - loss: 0.2572 -
mean_squared_error: 0.2572 - val_loss: 0.5339 - val_mean_squared_error: 0.5339
Epoch 78/100

130/130 [=====] - 0s 3ms/step - loss: 0.2381 -
mean_squared_error: 0.2381 - val_loss: 0.5302 - val_mean_squared_error: 0.5302
Epoch 79/100

130/130 [=====] - 0s 3ms/step - loss: 0.2293 -
mean_squared_error: 0.2293 - val_loss: 0.5112 - val_mean_squared_error: 0.5112
Epoch 80/100

130/130 [=====] - 0s 2ms/step - loss: 0.2370 -
mean_squared_error: 0.2370 - val_loss: 0.5078 - val_mean_squared_error: 0.5078
Epoch 81/100

130/130 [=====] - 0s 2ms/step - loss: 0.2315 -
mean_squared_error: 0.2315 - val_loss: 0.4973 - val_mean_squared_error: 0.4973
Epoch 82/100

130/130 [=====] - 0s 3ms/step - loss: 0.2296 -
mean_squared_error: 0.2296 - val_loss: 0.5417 - val_mean_squared_error: 0.5417
Epoch 83/100
130/130 [=====] - 0s 2ms/step - loss: 0.2241 -
mean_squared_error: 0.2241 - val_loss: 0.5404 - val_mean_squared_error: 0.5404
Epoch 84/100
130/130 [=====] - 0s 2ms/step - loss: 0.2262 -
mean_squared_error: 0.2262 - val_loss: 0.5882 - val_mean_squared_error: 0.5882
Epoch 85/100
130/130 [=====] - 0s 2ms/step - loss: 0.2276 -
mean_squared_error: 0.2276 - val_loss: 0.5403 - val_mean_squared_error: 0.5403
Epoch 86/100
130/130 [=====] - 0s 2ms/step - loss: 0.2213 -
mean_squared_error: 0.2213 - val_loss: 0.5317 - val_mean_squared_error: 0.5317
Epoch 87/100
130/130 [=====] - 0s 2ms/step - loss: 0.2154 -
mean_squared_error: 0.2154 - val_loss: 0.5465 - val_mean_squared_error: 0.5465
Epoch 88/100
130/130 [=====] - 0s 2ms/step - loss: 0.2271 -
mean_squared_error: 0.2271 - val_loss: 0.5295 - val_mean_squared_error: 0.5295
Epoch 89/100
130/130 [=====] - 0s 2ms/step - loss: 0.2131 -
mean_squared_error: 0.2131 - val_loss: 0.5705 - val_mean_squared_error: 0.5705
Epoch 90/100
130/130 [=====] - 0s 3ms/step - loss: 0.2260 -
mean_squared_error: 0.2260 - val_loss: 0.5359 - val_mean_squared_error: 0.5359
Epoch 91/100
130/130 [=====] - 0s 2ms/step - loss: 0.2154 -
mean_squared_error: 0.2154 - val_loss: 0.5451 - val_mean_squared_error: 0.5451
Epoch 92/100
130/130 [=====] - 0s 2ms/step - loss: 0.2021 -
mean_squared_error: 0.2021 - val_loss: 0.5389 - val_mean_squared_error: 0.5389
Epoch 93/100
130/130 [=====] - 0s 3ms/step - loss: 0.2029 -
mean_squared_error: 0.2029 - val_loss: 0.5346 - val_mean_squared_error: 0.5346
Epoch 94/100
130/130 [=====] - 0s 2ms/step - loss: 0.2106 -
mean_squared_error: 0.2106 - val_loss: 0.5739 - val_mean_squared_error: 0.5739
Epoch 95/100
130/130 [=====] - 0s 2ms/step - loss: 0.2074 -
mean_squared_error: 0.2074 - val_loss: 0.5516 - val_mean_squared_error: 0.5516
Epoch 96/100
130/130 [=====] - 0s 3ms/step - loss: 0.1997 -
mean_squared_error: 0.1997 - val_loss: 0.5699 - val_mean_squared_error: 0.5699
Epoch 97/100
130/130 [=====] - 0s 2ms/step - loss: 0.2046 -
mean_squared_error: 0.2046 - val_loss: 0.5623 - val_mean_squared_error: 0.5623
Epoch 98/100


```

130/130 [=====] - 0s 2ms/step - loss: 0.1976 -
mean_squared_error: 0.1976 - val_loss: 0.5765 - val_mean_squared_error: 0.5765
Epoch 99/100
130/130 [=====] - 0s 2ms/step - loss: 0.2179 -
mean_squared_error: 0.2179 - val_loss: 0.5628 - val_mean_squared_error: 0.5628
Epoch 100/100
130/130 [=====] - 0s 2ms/step - loss: 0.2009 -
mean_squared_error: 0.2009 - val_loss: 0.5468 - val_mean_squared_error: 0.5468

```

```

[42]: # Make predictions on the testing set
y_pred = ann_model.predict(X_test)

```

```

41/41 [=====] - 0s 1ms/step

```

```

[43]: # Calculate RMSE, MAE, MSE, and R2 score
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

```

```

[44]: print('RMSE:', rmse)
print('MAE:', mae)
print('MSE:', mse)
print('R-squared Score:', r2)

```

```

RMSE: 0.7477269652378632
MAE: 0.5659084646518414
MSE: 0.5590956145438246
R-squared Score: 0.24297870391632415

```

```

[45]: # separate the data and Label
x = wine_df.drop('quality',axis=1)
Y = wine_df['quality'].apply(lambda y_value: 1 if y_value>=7 else 0)

```

```

[46]: X_train, X_test, y_train, y_test = train_test_split(x, Y, test_size=0.2,
↳random_state=42)

```

0.5 4. Logistic Regression

```

[47]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

```

```

[48]: sc = StandardScaler()
X_train_std = sc.fit_transform(X_train)
X_test_std = sc.transform(X_test)

```

```
[49]: # create and fit the logistic regression model
lr = LogisticRegression(random_state=42)
lr.fit(X_train_std, y_train)
```

```
[49]: LogisticRegression(random_state=42)
```

```
[50]: # make predictions on the testing set
y_pred_lr = lr.predict(X_test_std)

# calculate the accuracy score and confusion matrix
acc_lr = accuracy_score(y_test, y_pred_lr)
cm_lr = confusion_matrix(y_test, y_pred_lr)
print('Accuracy Score (Logistic Regression):', acc_lr)
print('Confusion Matrix (Logistic Regression):\n', cm_lr)
```

```
Accuracy Score (Logistic Regression): 0.8246153846153846
Confusion Matrix (Logistic Regression):
[[1004  44]
 [ 184  68]]
```

0.6 5. Random Forest model

```
[51]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
[52]: # Create and train the random forest model
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
```

```
[52]: RandomForestClassifier()
```

```
[53]: # Evaluate the random forest model
rf_pred = rf_model.predict(X_test)
rf_acc = accuracy_score(y_test, rf_pred)
print("Random Forest accuracy: ", rf_acc)
```

```
Random Forest accuracy: 0.8861538461538462
```

```
[54]: # Create confusion matrix
y_true = y_test
y_pred = rf_pred
conf_mat = confusion_matrix(y_true, y_pred)
print(conf_mat)
```

```
[[1007  41]
 [ 107 145]]
```

0.7 6. Decision Tree

```
[55]: from sklearn.tree import DecisionTreeRegressor
      from sklearn.metrics import mean_squared_error
      import matplotlib.pyplot as plt
```

```
[56]: # Create a decision tree regressor with a maximum depth of 3
      tree_model = DecisionTreeRegressor(max_depth=1)
```

```
[57]: # Fit the model on the training data
      tree_model.fit(X_train, y_train)
```

```
[57]: DecisionTreeRegressor(max_depth=1)
```

```
[58]: # Make predictions on the test data
      dt_pred = tree_model.predict(X_test)
```

```
[59]: # Evaluate the model
      mse = mean_squared_error(y_test, dt_pred)
      print("Decision tree MSE: ", mse)
```

Decision tree MSE: 0.13583711324969888

0.8 LSTM Model

```
[60]: from sklearn.preprocessing import StandardScaler
      from keras.models import Sequential
      from keras.layers import Dense, LSTM, Dropout
```

```
[61]: X = wine.iloc[:, :-1].values
      y = wine.iloc[:, -1].values
```

```
[62]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[63]: sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[64]: X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
      X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
      # reshaping for LSTM model
```

```
[65]: model = Sequential()
      model.add(LSTM(50, return_sequences=True, input_shape=(X_train.shape[1], 1)))
      model.add(LSTM(50))
      model.add(Dense(1))
```

```
[66]: model.compile(optimizer='adam', loss='mean_squared_error')
```

```
[67]: model.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_test,   
↪y_test))
```

```
Epoch 1/100
163/163 [=====] - 6s 14ms/step - loss: 5.0015 -
val_loss: 0.6893
Epoch 2/100
163/163 [=====] - 1s 9ms/step - loss: 0.6909 -
val_loss: 0.6364
Epoch 3/100
163/163 [=====] - 1s 9ms/step - loss: 0.6278 -
val_loss: 0.5999
Epoch 4/100
163/163 [=====] - 2s 9ms/step - loss: 0.6144 -
val_loss: 0.6202
Epoch 5/100
163/163 [=====] - 2s 9ms/step - loss: 0.6094 -
val_loss: 0.5884
Epoch 6/100
163/163 [=====] - 2s 9ms/step - loss: 0.6027 -
val_loss: 0.5918
Epoch 7/100
163/163 [=====] - 2s 9ms/step - loss: 0.5995 -
val_loss: 0.5829
Epoch 8/100
163/163 [=====] - 2s 10ms/step - loss: 0.5935 -
val_loss: 0.5743
Epoch 9/100
163/163 [=====] - 2s 9ms/step - loss: 0.5942 -
val_loss: 0.5807
Epoch 10/100
163/163 [=====] - 2s 9ms/step - loss: 0.5865 -
val_loss: 0.5720
Epoch 11/100
163/163 [=====] - 2s 9ms/step - loss: 0.5804 -
val_loss: 0.5550
Epoch 12/100
163/163 [=====] - 2s 9ms/step - loss: 0.5765 -
val_loss: 0.6659
Epoch 13/100
163/163 [=====] - 2s 9ms/step - loss: 0.5790 -
val_loss: 0.5709
Epoch 14/100
163/163 [=====] - 2s 9ms/step - loss: 0.5725 -
val_loss: 0.5792
Epoch 15/100
163/163 [=====] - 2s 9ms/step - loss: 0.5662 -
val_loss: 0.5473
```

Epoch 16/100
163/163 [=====] - 2s 9ms/step - loss: 0.5657 -
val_loss: 0.5391
Epoch 17/100
163/163 [=====] - 2s 9ms/step - loss: 0.5587 -
val_loss: 0.5401
Epoch 18/100
163/163 [=====] - 2s 9ms/step - loss: 0.5599 -
val_loss: 0.5396
Epoch 19/100
163/163 [=====] - 2s 9ms/step - loss: 0.5520 -
val_loss: 0.5400
Epoch 20/100
163/163 [=====] - 2s 9ms/step - loss: 0.5581 -
val_loss: 0.5334
Epoch 21/100
163/163 [=====] - 2s 9ms/step - loss: 0.5549 -
val_loss: 0.5298
Epoch 22/100
163/163 [=====] - 2s 9ms/step - loss: 0.5509 -
val_loss: 0.5328
Epoch 23/100
163/163 [=====] - 2s 9ms/step - loss: 0.5451 -
val_loss: 0.5286
Epoch 24/100
163/163 [=====] - 2s 9ms/step - loss: 0.5478 -
val_loss: 0.5266
Epoch 25/100
163/163 [=====] - 2s 9ms/step - loss: 0.5434 -
val_loss: 0.5303
Epoch 26/100
163/163 [=====] - 2s 9ms/step - loss: 0.5442 -
val_loss: 0.5183
Epoch 27/100
163/163 [=====] - 2s 9ms/step - loss: 0.5377 -
val_loss: 0.5503
Epoch 28/100
163/163 [=====] - 2s 9ms/step - loss: 0.5386 -
val_loss: 0.5178
Epoch 29/100
163/163 [=====] - 2s 9ms/step - loss: 0.5351 -
val_loss: 0.5220
Epoch 30/100
163/163 [=====] - 2s 9ms/step - loss: 0.5353 -
val_loss: 0.5334
Epoch 31/100
163/163 [=====] - 2s 9ms/step - loss: 0.5326 -
val_loss: 0.5157

Epoch 32/100
163/163 [=====] - 2s 9ms/step - loss: 0.5262 -
val_loss: 0.5191
Epoch 33/100
163/163 [=====] - 2s 9ms/step - loss: 0.5268 -
val_loss: 0.5166
Epoch 34/100
163/163 [=====] - 2s 9ms/step - loss: 0.5265 -
val_loss: 0.5188
Epoch 35/100
163/163 [=====] - 2s 10ms/step - loss: 0.5225 -
val_loss: 0.5467
Epoch 36/100
163/163 [=====] - 2s 9ms/step - loss: 0.5239 -
val_loss: 0.5194
Epoch 37/100
163/163 [=====] - 2s 9ms/step - loss: 0.5196 -
val_loss: 0.5249
Epoch 38/100
163/163 [=====] - 2s 9ms/step - loss: 0.5165 -
val_loss: 0.5272
Epoch 39/100
163/163 [=====] - 2s 9ms/step - loss: 0.5119 -
val_loss: 0.5111
Epoch 40/100
163/163 [=====] - 2s 9ms/step - loss: 0.5107 -
val_loss: 0.5139
Epoch 41/100
163/163 [=====] - 2s 10ms/step - loss: 0.5105 -
val_loss: 0.5167
Epoch 42/100
163/163 [=====] - 2s 9ms/step - loss: 0.5130 -
val_loss: 0.5200
Epoch 43/100
163/163 [=====] - 2s 10ms/step - loss: 0.5063 -
val_loss: 0.5036
Epoch 44/100
163/163 [=====] - 2s 9ms/step - loss: 0.5090 -
val_loss: 0.5062
Epoch 45/100
163/163 [=====] - 2s 10ms/step - loss: 0.5073 -
val_loss: 0.5052
Epoch 46/100
163/163 [=====] - 2s 9ms/step - loss: 0.4995 -
val_loss: 0.5070
Epoch 47/100
163/163 [=====] - 2s 9ms/step - loss: 0.4991 -
val_loss: 0.5002

Epoch 48/100
163/163 [=====] - 2s 9ms/step - loss: 0.4947 -
val_loss: 0.5184
Epoch 49/100
163/163 [=====] - 2s 9ms/step - loss: 0.4921 -
val_loss: 0.5108
Epoch 50/100
163/163 [=====] - 2s 9ms/step - loss: 0.4924 -
val_loss: 0.4884
Epoch 51/100
163/163 [=====] - 2s 9ms/step - loss: 0.4887 -
val_loss: 0.4932
Epoch 52/100
163/163 [=====] - 2s 9ms/step - loss: 0.4875 -
val_loss: 0.4911
Epoch 53/100
163/163 [=====] - 2s 9ms/step - loss: 0.4844 -
val_loss: 0.4874
Epoch 54/100
163/163 [=====] - 2s 10ms/step - loss: 0.4813 -
val_loss: 0.4889
Epoch 55/100
163/163 [=====] - 2s 10ms/step - loss: 0.4794 -
val_loss: 0.4921
Epoch 56/100
163/163 [=====] - 2s 10ms/step - loss: 0.4800 -
val_loss: 0.4922
Epoch 57/100
163/163 [=====] - 2s 10ms/step - loss: 0.4715 -
val_loss: 0.4906
Epoch 58/100
163/163 [=====] - 2s 10ms/step - loss: 0.4716 -
val_loss: 0.4887
Epoch 59/100
163/163 [=====] - 2s 10ms/step - loss: 0.4697 -
val_loss: 0.4952
Epoch 60/100
163/163 [=====] - 2s 10ms/step - loss: 0.4629 -
val_loss: 0.4847
Epoch 61/100
163/163 [=====] - 2s 10ms/step - loss: 0.4667 -
val_loss: 0.5008
Epoch 62/100
163/163 [=====] - 2s 10ms/step - loss: 0.4609 -
val_loss: 0.4781
Epoch 63/100
163/163 [=====] - 2s 10ms/step - loss: 0.4590 -
val_loss: 0.4877

Epoch 64/100
163/163 [=====] - 2s 10ms/step - loss: 0.4525 -
val_loss: 0.4825
Epoch 65/100
163/163 [=====] - 2s 10ms/step - loss: 0.4507 -
val_loss: 0.4956
Epoch 66/100
163/163 [=====] - 2s 10ms/step - loss: 0.4485 -
val_loss: 0.4928
Epoch 67/100
163/163 [=====] - 2s 10ms/step - loss: 0.4490 -
val_loss: 0.4811
Epoch 68/100
163/163 [=====] - 2s 10ms/step - loss: 0.4438 -
val_loss: 0.4783
Epoch 69/100
163/163 [=====] - 2s 10ms/step - loss: 0.4381 -
val_loss: 0.4967
Epoch 70/100
163/163 [=====] - 2s 10ms/step - loss: 0.4392 -
val_loss: 0.4900
Epoch 71/100
163/163 [=====] - 2s 10ms/step - loss: 0.4356 -
val_loss: 0.4852
Epoch 72/100
163/163 [=====] - 2s 10ms/step - loss: 0.4321 -
val_loss: 0.5009
Epoch 73/100
163/163 [=====] - 2s 10ms/step - loss: 0.4298 -
val_loss: 0.4862
Epoch 74/100
163/163 [=====] - 2s 10ms/step - loss: 0.4248 -
val_loss: 0.4959
Epoch 75/100
163/163 [=====] - 2s 10ms/step - loss: 0.4240 -
val_loss: 0.4910
Epoch 76/100
163/163 [=====] - 2s 10ms/step - loss: 0.4173 -
val_loss: 0.4885
Epoch 77/100
163/163 [=====] - 2s 10ms/step - loss: 0.4157 -
val_loss: 0.4870
Epoch 78/100
163/163 [=====] - 2s 10ms/step - loss: 0.4143 -
val_loss: 0.4900
Epoch 79/100
163/163 [=====] - 2s 10ms/step - loss: 0.4092 -
val_loss: 0.4871

Epoch 80/100
163/163 [=====] - 2s 10ms/step - loss: 0.4039 -
val_loss: 0.5206
Epoch 81/100
163/163 [=====] - 2s 10ms/step - loss: 0.4016 -
val_loss: 0.5032
Epoch 82/100
163/163 [=====] - 2s 9ms/step - loss: 0.3996 -
val_loss: 0.4990
Epoch 83/100
163/163 [=====] - 2s 9ms/step - loss: 0.3938 -
val_loss: 0.5020
Epoch 84/100
163/163 [=====] - 1s 9ms/step - loss: 0.3913 -
val_loss: 0.4963
Epoch 85/100
163/163 [=====] - 1s 9ms/step - loss: 0.3850 -
val_loss: 0.4938
Epoch 86/100
163/163 [=====] - 1s 9ms/step - loss: 0.3881 -
val_loss: 0.4916
Epoch 87/100
163/163 [=====] - 1s 9ms/step - loss: 0.3768 -
val_loss: 0.5014
Epoch 88/100
163/163 [=====] - 2s 9ms/step - loss: 0.3788 -
val_loss: 0.4978
Epoch 89/100
163/163 [=====] - 1s 9ms/step - loss: 0.3704 -
val_loss: 0.4857
Epoch 90/100
163/163 [=====] - 1s 9ms/step - loss: 0.3678 -
val_loss: 0.5123
Epoch 91/100
163/163 [=====] - 1s 9ms/step - loss: 0.3643 -
val_loss: 0.5088
Epoch 92/100
163/163 [=====] - 1s 9ms/step - loss: 0.3602 -
val_loss: 0.5078
Epoch 93/100
163/163 [=====] - 1s 9ms/step - loss: 0.3548 -
val_loss: 0.5048
Epoch 94/100
163/163 [=====] - 1s 9ms/step - loss: 0.3529 -
val_loss: 0.5147
Epoch 95/100
163/163 [=====] - 2s 9ms/step - loss: 0.3467 -
val_loss: 0.5100

```

Epoch 96/100
163/163 [=====] - 2s 10ms/step - loss: 0.3405 -
val_loss: 0.5173
Epoch 97/100
163/163 [=====] - 2s 10ms/step - loss: 0.3411 -
val_loss: 0.5060
Epoch 98/100
163/163 [=====] - 2s 10ms/step - loss: 0.3360 -
val_loss: 0.5060
Epoch 99/100
163/163 [=====] - 2s 9ms/step - loss: 0.3306 -
val_loss: 0.5179
Epoch 100/100
163/163 [=====] - 1s 9ms/step - loss: 0.3279 -
val_loss: 0.5073

```

[67]: <keras.callbacks.History at 0x243360ba830>

```
[68]: loss = model.evaluate(X_test, y_test)
      print('Test Loss:', loss)
```

```

41/41 [=====] - 0s 4ms/step - loss: 0.5073
Test Loss: 0.5073338150978088

```

0.9 Multilayer Perceptron (MLP)

```
[69]: from sklearn.neural_network import MLPRegressor
```

```
[70]: # split the data into train and test sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

```
[71]: # standardize the data
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[72]: mlp = MLPRegressor(hidden_layer_sizes=(50, 50), activation='relu',
      ↪solver='adam', max_iter=500, random_state=42)
      mlp.fit(X_train, y_train)
```

[72]: MLPRegressor(hidden_layer_sizes=(50, 50), max_iter=500, random_state=42)

```
[73]: # evaluate the model on the test set
      y_pred = mlp.predict(X_test)
      mse = mean_squared_error(y_test, y_pred)
      print("Test MSE: ", mse)
```

Test MSE: 0.4930566425509396

0.10 Prediction

```
[74]: input_data = (14,0.5,0.36,6.1,0.071,17.0,102.0,0.9978,3.35,0.8,10.5)
```

```
[75]: # changing the input data to a numpy array  
input_data_as_numpy_array = np.asarray(input_data)
```

```
[76]: # reshape the data as we are predicting the label for only one instance  
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
[77]: prediction = rf_model.predict(input_data_reshaped)  
print(prediction)  
  
if (prediction[0]==1):  
    print('Good Quality Wine')  
else:  
    print('Bad Quality Wine')
```

```
[0]
```

```
Bad Quality Wine
```